

DANIEL SPIRIDONOV

Machine Learning Engineer · Moscow, Russia

Phone: +7 (936) 505-12-45 | Telegram: [@daspiridonov](https://t.me/daspiridonov) | sprv-4@yandex.ru
linkedin.com/in/daniel-spiridonov | github.com/cop1cat

Summary

ML Engineer with 1.5 years of production experience. Interested in engineering ML/LLM systems: agentic pipelines, inference infrastructure, MLOps. 4th-year student at HSE University.

Work Experience

RT Information Security — Machine Learning Engineer

Dec 2024 - Present

- Built a CatBoost alert classifier: training with PCA and custom embeddings of domain text, multi-threaded online inference with hot-swap weights. 99% recall on TP, 20–30% on FP — drops a portion of false alerts
- Assembled a pipeline for text-exception generation: regex pre-filter, Levenshtein deduplication, an LLM stage with self-validation via LangGraph, inference on self-hosted vLLM. Rules used to be written by hand, now an analyst only reviews the generated ones — covers ~70% of cases
- Shipped a production LLM-agent service (Python, LangGraph) for incident triage: state graph, tools, structured output, inter-step context compression, observability via Langfuse. Produces an incident breakdown used by analysts as a reference before escalation
- Deployed an anomaly detector over ~100 event sources with auto method selection: z-score for source-down detection, time-series decomposition with seasonality and a quantile threshold for traffic anomalies. Runs every 10 minutes; results go to RabbitMQ and PostgreSQL
- Refactored an incident-grouping service (finds a similar incident and sets a parent/child link): decoupled components, swapping retrieval models and comparing them became easier. Covers ~15% of all groupings
- Wrote a pip package for feature and inference pipelines with Optuna hyperparameter search — a team standard, speeds up new-project bootstrapping several-fold
- Set up the team's MLOps infrastructure: CI/CD for ML services with docker-compose deployment, MinIO for dataset and model-artifact versioning, team-wide code-quality standards (ruff, mypy, docstrings)
- Participated in technical interviews of candidates and onboarding of new team members

Technical Skills

- **Programming Languages:** Python, SQL
- **Frameworks & Libraries:** pandas, numpy, scikit-learn, CatBoost, XGBoost, PyTorch, FastAPI, LangGraph, vLLM, Langfuse
- **MLOps & Infrastructure:** Docker, GitLab CI/CD, Airflow, Prefect, MLflow, MinIO, PostgreSQL, Grafana
- **Dev Tools:** git, ruff, mypy, uv
- **Spoken Languages:** English (B2), Russian (native)

Pet Projects

- [fluxio](#) — pipeline framework for Python with auto-parallelism, crash recovery, and streaming
- [catboost-utils](#) — CatBoost wrapper with readable errors, data validation, custom losses, sklearn integration
- [snowcite](#) — MCP server for literature review: search across arXiv/Semantic Scholar/OpenAlex/PubMed, PDF via LaTeX/Typst
- [vaultly](#) — Pydantic-native secrets manager with lazy loading, TTL caching, and log masking

Education

HSE University — Specialist Degree (5-year, equivalent to MSc) in Computer and Information Systems Security

2022 — 2028 (*expected*)